

Yu-Neng (Allen) Chuang

☎ (979) 334-0045 ◇ ✉ yc146@rice.edu ◇ 🌐 Homepage: <https://ynchuang.github.io>

RESEARCH INTERESTS & SKILLS

I aim to build reliable and efficient LLM systems through techniques such as efficient long-context reasoning, mechanism interpretation, uncertainty-guided routing, and multimodal agent systems.

- **LLM Agent & Post-training:** Multimodal Agents, Reasoning, Routing, Long-context Efficiency
- **Trustworthy LLMs:** Interpretation, Uncertainty, Safety, Evaluation

EDUCATION

Rice University (Co-advised by Dr. Xia “Ben” Hu & Dr. Vladimir Braverman) Houston, TX
Ph.D. in Computer Science Aug. 2021 - Present

National ChengChi University Taipei, Taiwan
Master of Science in Computer Science Feb. 2018 - Jun. 2020

National ChengChi University Taipei, Taiwan
Bachelor of Science in Mathematical Science Aug. 2013 - Jul. 2017

EXPERIENCE

Google DeepMind @ Gemini Team New York City, NY
Research Intern May 2025 - Aug. 2025

- Developed multi-modal agentic LLMs leveraging the Model Context Protocol (MCP) with reasoning for dynamic tool planning and usage in complex audio tasks, reducing system failure rates by 28% and achieving hallucination-free tool calling.

Apple Inc. @ Siri Team Cupertino, CA
Research Intern May 2024 - Aug. 2024

- Developed uncertainty quantification algorithm for on-device LLM routing purpose based on confidence token prediction, deducing 52% of inference latency while maintaining the performance

Samsung Research America @ LLM Team Mountain View, CA
Research Intern May 2023 - Aug. 2023

- Developed an efficient algorithm of hard prompt compression on large language models with LLM post-training techniques, Reinforcement learning with human feedback (RLHF) method, deducing 80% of LLM API usage cost and 20% of latency of white box LLMs

Rice University Houston, TX
Graduate Research Assistant Aug. 2021 - Present

- Developed fine-tuning and alignment algorithms to address uncertainty, enhance explainability, and improve safety in LLMs
- Investigated multi-agent systems to enhance LLM reasoning, planning, and inference routing
- Built efficient framework for KV-cache lossy compression for lower inference throughput and budgets

Carnegie Mellon University and SMU @ Living Analytics Research Centre Singapore
Research Assistant Jan. 2020 - Apr. 2020

- Built a ranking method for a personalized job recommendation system with TB-scaled user data in Singapore, which outperformed other state-of-the-art ranking methods by 10.2%

KKBOX Co, Ltd. Taipei, Taiwan
Data Scientist Intern Sep. 2019 - Jun. 2020

- Developed a ranking algorithm on a streaming dataset of nearly 1.5 million users to enhance the recommendation system by 15.3% compared with the prior internal recommendation systems

SELECTED PUBLICATIONS

Conference and Journal Publications

- [ICML' 25] **Y.N. Chuang**, H. Zhou, P. Sarma, P. Gopalan, J. Boccio, S. Bolouki, and X. Hu. "Learning to Route LLMs with Confidence Tokens" *International Conference on Machine Learning*
- [NeurIPS' 25] **Y.N. Chuang***, F. Lou*, et. al., V. Braverman, V. Chaudhary, and X. Hu. "AutoL2S: Auto Long-Short Reasoning for Efficient Large Language Models" *Annual Conference on Neural Information Processing Systems Workshop on Efficient Reasoning*
- [EMNLP' 24] **Y.N. Chuang***, G. Wang*, R. Tang, S. Zhong, J. Yuan, H. Jin, Z. Liu, V. Chaudhary, S. Xu, J. Caverlee, and X. Hu. "Taylor Unswift: Secured Weight Release for Large Language Models via Taylor Expansion" *Annual Conference of the North American Chapter of the ACL*
- [NAACL' 24] **Y.N. Chuang**, T. Xing, C.Y. Chang, Z. Liu, X. Chen, and X. Hu. "Learning to Compress Prompt in Natural Language Formats" *Annual Conference of the North American Chapter of the ACL*
- [NAACL' 24 Finding] **Y.N. Chuang***, R. Tang*, and X. Hu. "Secure Your Model: A Simple but Effective Key Prompt Protection Mechanism for Large Language Models" *Finding of Annual Conference of the North American Chapter of the ACL*
- [ICLR' 23] **Y.N. Chuang***, G. Wang*, F. Yang, Q. Zhou, P. Tripathi, X. Cai and X. Hu. "CoRTX: Contrastive Learning for Real-time Explanations" *International Conference on Learning Representations*
- [ICML' 22 Spotlight] **Y.N. Chuang***, G. Wang*, M. Du, F. Yang, Q. Zhou, P. Tripathi, X. Cai and X. Hu. "Accelerating Shapley Explanation via Contributive Cooperator Selection" *International Conference on Machine Learning*
- [JBI] **Y.N. Chuang**, R. Tang, X. Jiang, and X. Hu. "SPeC: A Soft Prompt-Based Calibration on Performance Variability of Large Language Model in Clinical Notes Summarization" *Journal of Biomedical Informatics*
- [TKDD] **Y.N. Chuang**, Cheng-Te Li. "Privacy-Preserving Representation Learning with Gradient Obfuscation against Attribute Inference for Recommendation" *ACM Transactions on Knowledge Discovery from Data*
- [TKDD] **Y.N. Chuang**, K.H. Lai, R. Tang, M. Du, C.Y. Chang, N. Zou, and X. Hu. "Mitigating Relational Bias on Knowledge Graphs" *ACM Transactions on Knowledge Discovery from Data*
- [CIKM' 23 Best Paper Runner-up] **Y.N. Chuang**, G. Wang et al., and X. Hu "DiscoverPath: A Knowledge Refinement and Retrieval System for Interdisciplinarity on Biomedical Research" *ACM International Conference on Information and Knowledge Management (CIKM'23 Best Demo Paper Honorable Mention)*
- [CIKM' 20] **Y.N. Chuang***, C.M. Chen*, C.J. Wang, M.F. Tsai, Y. Fang, and E.P. Lim. "TPR: Text-aware Preference Ranking for Recommender Systems" *ACM International Conference on Information and Knowledge Management*
- [UAI' 20 Oral] **Y.N. Chuang***, C.J. Wang*, C.M. Chen, and M.F. Tsai. "Skewness Ranking Optimization for Personalized Recommendation" *Conference on Uncertainty in Artificial Intelligence (Oral)*
- [CACM] R. Tang, **Y.N. Chuang**, and X. Hu. "The Science of LLM-generated Text Detection" *The Communications of the ACM (CACM April Cover)*
- [TMLR] Y. Sui, **Y.N. Chuang**, G. Wang, et. al., and X. Hu. "Stop Overthinking: A Survey on Efficient Reasoning for Large Language Models" *Transactions on Machine Learning Research*
- [NeurIPS' 25] Z. Xu*, G. Wang*, **Y.N. Chuang**, G. Zheng, A. Szalay, Z. Liu, V. Braverman "DTS: Enhancing Large Reasoning Models via Decoding Tree Sketching" *Annual Conference on Neural Information Processing Systems Workshop on Efficient Reasoning*
- [KDD' 25] C.Y. Chang, **Y.N. Chuang**, Z. Jiang, K.H. Lai, A. Jiang, N. Zou. "CODA: Temporal Domain Generalization via Concept Drift Simulator" *International Conference on Knowledge Discovery and Data Mining*

[ACL' 25] M. Zhong, G. Wang, **Y.N. Chuang**, N. Zou. "Quantized Can Still Be Calibrated: A Unified Framework to Calibration in Quantized Large Language Models" *Annual Meeting of the Association for Computational Linguistics*

[ACL' 25 Finding] J. Zhang, J. Yuan, A. Wen, H. Le, **Y.N. Chuang**, S. Choi, R. Chen, X. Hu. "ReasonerRank: Redefining Language Model Evaluation with Ground-Truth-Free Ranking Frameworks" *Finding of Annual Meeting of the Association for Computational Linguistics*

[EMNLP' 25 Finding] L. Zhang, **Y.N. Chuang**, G. Wang, R. Tang, X. Cai, R. Shenoy, and X. Hu "A Decoupled Multi-Agent Framework for Complex Text Style Transfer"

[EMNLP' 25 Finding] Z. Xu, G. Wang, G. Zheng, **Y.N. Chuang**, A. Szalay, X. Hu, and V. Braverman "Self-Ensemble: Mitigating Confidence Distortion for Large Language Models"

[NAACL' 25 Finding] Y. Wang*, J. Yuan*, **Y.N. Chuang**, et al, and X. Hu, "DHP Benchmark: Are LLMs Good NLG Evaluators?" *Finding of Annual Conference of the North American Chapter of the ACL*

[ICML' 24] G. Wang, **Y.N. Chuang**, F. Yang, M. Du, C.Y. Chang, et al., and X. Cai, and X. Hu. "TVE: Learning Meta-attribution for Transferable Vision Explainer" *International Conference on Machine Learning*

[EMNLP' 24 Finding] J. Yuan*, H. Liu*, S. Zhong*, **Y.N. Chuang**, et al., and X. Hu. "KV Cache Compression, But What Must We Give in Return? A Comprehensive Benchmark of Long Context Capable Approaches" *Finding of Empirical Methods in Natural Language Processing*

Preprints

[Submitted ARR] **Y.N. Chuang***, L. Yu, G. Wang, et. al., V. Braverman, and X. Hu. "Confident or Seek Stronger: Exploring Uncertainty-Based On-device LLM Routing From Benchmarking to Generalization"

[Submitted ARR] **Y.N. Chuang**, G. Wang, C.Y. Chang, R. Tang, S. Zhong, F. Yang, M. Du, X. Cai, V. Braverman, and X. Hu "FaithLM: Towards Faithful Explanations for Large Language Models"

OPEN SOURCE PROJECTS

DiscoverPath: A Knowledge Refinement and Retrieval System for Interdisciplinarity on Biomedical Research (CIKM'23 Best Demo Paper Honorable Mention)

- *Project Leader.* Designed a KG-based retrieval system designed for biomedical research that aims to assist biomedical researchers in dynamically refining their queries and effectively retrieving articles.

LTSM-bundle: Large Time Series Models Training and Benchmark Library

- *Project Leader.* Designed the package architectures with CI/CD pipeline for large-scale time series data.

SMORe: Modularize Graph Embedding for Recommendation

- *Developer.* Constructed a large-scale network embedding library for recommendation systems on online streaming services which was developed under C++ with multi-thread processing techniques

HONORS AND AWARDS

- Rice D2K Fellowship, Rice University	Sep. 2025
- Study Abroad Fellowship, Ministry of Education, Taiwan	May. 2025
- Ken Kennedy Institute Fellowship, Rice University	Nov. 2024
- Doctoral Forum Travel Award, SDM' 24	Mar. 2024
- CIKM 2023 Best Demo Paper Honorable Mention	Oct. 2023
- 4th Place at ACM RecSys Challenge	Sep. 2020

PROFESSIONAL SERVICES

Reviewer (Since 2020): NeurIPS, ICLR, ICML, ACL, EMNLP, NAACL, AAAI, IJCAI, KDD, WSDM, CIKM, IEEE TPAMI, IEEE TAI, IEEE TIST, IEEE ICHI